

Dative_Alternation_Analysis_Using_Logistic_Regression

June 2, 2023

0.1 A new application of logistic regression: the dative alternation

The work you need to do for this pset involves applying logistic regression to a new case, the **dative alternation**, which we studied in a previous pset. We will use the **dative** dataset from Bresnan et al. (2007). First we load the dataset:

```
[ ]: # imports
import statsmodels.api as sm
import math, random
import pandas as pd
import numpy as np

[ ]: dat = pd.read_csv("https://gist.githubusercontent.com/scaperex/
↳278815a736401d36021aa9fe31b9a0cb/raw/
↳cf338a8cf745fa5820c4ea97af682d265bc1a34f/dative-alternation.csv")
dat
```

```
[ ]: Unnamed: 0 Speaker Modality Verb SemanticClass LengthOfRecipient \
0          1      NaN written feed t 1
1          2      NaN written give a 2
2          3      NaN written give a 1
3          4      NaN written give a 1
4          5      NaN written offer c 2
...      ...      ...      ...      ...      ...
3258      3258    S1190 spoken tell c 1
3259      3259    S1423 spoken give a 1
3260      3260    S1680 spoken give a 4
3261      3261    S1680 spoken give a 1
3262      3262    S1023 spoken pay a 1

AnimacyOfRec DefinOfRec PronomOfRec LengthOfTheme AnimacyOfTheme \
0 animate definite pronominal 14 inanimate
1 animate definite nonpronominal 3 inanimate
2 animate definite nonpronominal 13 inanimate
3 animate definite pronominal 5 inanimate
4 animate definite nonpronominal 3 inanimate
...      ...      ...      ...      ...
3258 animate definite pronominal 1 inanimate
```

3259	animate	definite	pronominal	9	inanimate
3260	animate	indefinite	nonpronominal	2	inanimate
3261	inanimate	definite	pronominal	2	inanimate
3262	animate	definite	pronominal	2	inanimate

	DefinOfTheme	PronomOfTheme	RealizationOfRecipient	AccessOfRec	\
0	indefinite	nonpronominal	NP	given	
1	indefinite	nonpronominal	NP	given	
2	definite	nonpronominal	NP	given	
3	indefinite	nonpronominal	NP	given	
4	definite	nonpronominal	NP	given	
...	
3258	definite	pronominal	NP	given	
3259	indefinite	nonpronominal	NP	given	
3260	definite	nonpronominal	PP	accessible	
3261	indefinite	nonpronominal	NP	given	
3262	indefinite	nonpronominal	NP	given	

	AccessOfTheme
0	new
1	new
2	new
3	new
4	new
...	...
3258	given
3259	accessible
3260	accessible
3261	accessible
3262	accessible

[3263 rows x 16 columns]

We see that it uses text values for some of the variables we are interested in (the response variable `RealizationOfRecipient`, and the variables expressing length and pronominality of theme and object). We create numeric versions of these variables, arbitrarily coding a double object outcome as 1 (“success”) and a prepositional dative outcome as 0.

```
[ ]: dat["Response"] = [1 if x == "NP" else 0 for x in dat["RealizationOfRecipient"]]
dat["RecPro"] = [1 if x == "pronominal" else 0 for x in dat["PronomOfRec"]]
dat["ThemePro"] = [1 if x == "pronominal" else 0 for x in dat["PronomOfTheme"]]
dat[["RealizationOfRecipient", "Response", "PronomOfRec", "RecPro", "ThemePro"]]
```

	RealizationOfRecipient	Response	PronomOfRec	RecPro	ThemePro
0	NP	1	pronominal	1	0
1	NP	1	nonpronominal	0	0
2	NP	1	nonpronominal	0	0

3	NP	1	pronominal	1	0
4	NP	1	nonpronominal	0	0
...
3258	NP	1	pronominal	1	1
3259	NP	1	pronominal	1	0
3260	PP	0	nonpronominal	0	0
3261	NP	1	pronominal	1	0
3262	NP	1	pronominal	1	0

[3263 rows x 5 columns]

```
[ ]: ## TODO: create numeric variables for PronomOfTheme
dat["ThemePro"] = [1 if x == "pronominal" else 0 for x in dat["PronomOfTheme"]]
dat[["ThemePro"]]
```

```
[ ]:      ThemePro
0          0
1          0
2          0
3          0
4          0
...
3258       1
3259       0
3260       0
3261       0
3262       0
```

[3263 rows x 1 columns]

To capture the possibility of an overall preference for one construction or the other, we add an “intercept” term to the logistic regression model, by creating a new **Dummy** variable in the data frame. We then fit a baseline model using only the intercept and find that there is an overall majority preference for the **DO** realization in this dataset (the intercept’s fitted weight is greater than 0). We also see that the intercept-only model simply recapitulates the sample mean.

```
[ ]: dat["Dummy"] = 1
x = dat[["Dummy"]]
y = dat[["Response"]]
model_0 = sm.GLM(y,x,family=sm.families.Binomial()) # first argument is
    ↪ response, second argument is predictor matrix, third argument says this is
    ↪ logistic regression
model_0 = model_0.fit()
print(model_0.summary())
print("Predicted proportion of DO outcomes based on fitted intercept-only model:
    ↪", round(np.mean(model_0.predict(x)),4))
```

```
print("Proportion of data with D0 outcome:", round(np.mean(y["Response"]),4)) #_
↳ same as model-predicted proportion
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:          Response    No. Observations:          3263
Model:                  GLM        Df Residuals:              3262
Model Family:          Binomial    Df Model:                0
Link Function:         Logit       Scale:                  1.0000
Method:                IRLS       Log-Likelihood:         -1870.5
Date:                  Sat, 13 May 2023    Deviance:              3741.1
Time:                  19:21:49    Pearson chi2:          3.26e+03
No. Iterations:        4          Pseudo R-squ. (CS):     -2.220e-16
Covariance Type:      nonrobust
=====
```

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
Dummy          1.0450      0.040     26.189      0.000      0.967      1.123
=====
```

Predicted proportion of D0 outcomes based on fitted intercept-only model: 0.7398

Proportion of data with D0 outcome: 0.7398

Task: In the below code boxes, complete the five parts of the problem specified in the pset PDF.

```
[ ]: def calc_accuracy(predicted_values, true_values):
    s = 0

    for pred_val, true_val in zip(predicted_values, true_values):

        if pred_val > 0.5:
            pred_val = 1

        else:
            pred_val = 0

        if pred_val == true_val:
            s += 1

    acc = s / len(true_values)
    return acc

def calc_loglikelihood(predicted_values, true_values):
    log_likelihood = 0

    for pred_val, true_val in zip(predicted_values, true_values):
        # we use log with base 2 according to the course's convention
```

```

    log_likelihood += (true_val * math.log2(pred_val)) + ((1 - true_val) * math.
↪log2(1 - pred_val))

    return log_likelihood

def eval_model(model, test_set, predictors):
    X_test = test_set[predictors]
    y_predicted, y = model.predict(X_test), test_set["Response"]

    model_accuracy = calc_accuracy(y_predicted, y)
    model_log_likelihood = calc_loglikelihood(y_predicted, y)

    print(f"model accuracy on the test set: {round(model_accuracy, 3)}")
    print(f"model log-likelihood on the test set: {round(model_log_likelihood,
↪3)}")

```

0.1.1 Part 1

implementation of 80/20 train/test random split of the “dative” dataset

```

[ ]: # setting "random_set" in order get consistent results
train_set = dat.sample(frac=0.8, random_state=0)
test_set = dat.drop(train_set.index.tolist())

# reset indicies
train_set.reset_index(drop=True, inplace=True)
test_set.reset_index(drop=True, inplace=True)

display(train_set)
display(test_set)

```

	Unnamed: 0	Speaker	Modality	Verb	SemanticClass	LengthOfRecipient	\
0	2464	S1421	spoken	give	t	1	
1	1515	S1151	spoken	send	t	1	
2	2756	S1462	spoken	teach	c	1	
3	655	NaN	written	give	a	8	
4	1998	S1259	spoken	tell	c	1	
...	
2605	2709	S1488	spoken	tell	c	1	
2606	1181	S1022	spoken	give	c	1	
2607	338	NaN	written	pay	t	1	
2608	47	NaN	written	offer	t	1	
2609	748	NaN	written	funnel	t	4	

	AnimacyOfRec	DefinOfRec	PronomOfRec	LengthOfTheme	AnimacyOfTheme	\
0	animate	definite	pronominal	2	inanimate	

1	animate	definite	pronominal	2	inanimate
2	animate	definite	pronominal	2	inanimate
3	inanimate	indefinite	nonpronominal	7	inanimate
4	animate	definite	pronominal	3	inanimate
...
2605	animate	definite	pronominal	3	inanimate
2606	animate	definite	pronominal	3	inanimate
2607	animate	indefinite	nonpronominal	6	inanimate
2608	animate	indefinite	nonpronominal	3	inanimate
2609	animate	definite	nonpronominal	1	inanimate

	DefinOfTheme	PronomOfTheme	RealizationOfRecipient	AccessOfRec	\
0	indefinite	nonpronominal	NP	given	
1	indefinite	pronominal	NP	given	
2	indefinite	nonpronominal	NP	given	
3	indefinite	nonpronominal	PP	new	
4	definite	nonpronominal	NP	given	
...	
2605	definite	pronominal	NP	given	
2606	definite	nonpronominal	NP	given	
2607	indefinite	nonpronominal	NP	accessible	
2608	indefinite	nonpronominal	NP	given	
2609	indefinite	nonpronominal	PP	new	

	AccessOfTheme	Response	RecPro	ThemePro	Dummy
0	new	1	1	0	1
1	accessible	1	1	1	1
2	new	1	1	0	1
3	new	0	0	0	1
4	given	1	1	0	1
...
2605	given	1	1	1	1
2606	accessible	1	1	0	1
2607	new	1	0	0	1
2608	new	1	0	0	1
2609	new	0	0	0	1

[2610 rows x 20 columns]

	Unnamed: 0	Speaker	Modality	Verb	SemanticClass	LengthOfRecipient	\
0	1	NaN	written	feed	t	1	
1	4	NaN	written	give	a	1	
2	8	NaN	written	bring	a	1	
3	22	NaN	written	give	a	1	
4	25	NaN	written	give	a	2	
..	
648	3240	S1697	spoken	give	c	1	
649	3243	S1621	spoken	give	a	1	

650	3244	S1621	spoken	pay	a	4
651	3251	S1382	spoken	tell	c	1
652	3257	S1701	spoken	send	t	1

	AnimacyOfRec	DefinOfRec	PronomOfRec	LengthOfTheme	AnimacyOfTheme	\
0	animate	definite	pronominal	14	inanimate	
1	animate	definite	pronominal	5	inanimate	
2	animate	definite	pronominal	1	inanimate	
3	animate	definite	nonpronominal	7	inanimate	
4	animate	indefinite	nonpronominal	10	inanimate	
..	
648	animate	definite	pronominal	3	inanimate	
649	animate	definite	pronominal	3	inanimate	
650	inanimate	definite	nonpronominal	10	inanimate	
651	animate	definite	pronominal	2	inanimate	
652	animate	definite	nonpronominal	4	inanimate	

	DefinOfTheme	PronomOfTheme	RealizationOfRecipient	AccessOfRec	\
0	indefinite	nonpronominal	NP	given	
1	indefinite	nonpronominal	NP	given	
2	indefinite	nonpronominal	NP	given	
3	indefinite	nonpronominal	NP	given	
4	definite	nonpronominal	NP	accessible	
..	
648	indefinite	nonpronominal	NP	given	
649	indefinite	nonpronominal	NP	given	
650	indefinite	nonpronominal	PP	accessible	
651	indefinite	nonpronominal	NP	given	
652	indefinite	nonpronominal	PP	new	

	AccessOfTheme	Response	RecPro	ThemePro	Dummy
0	new	1	1	0	1
1	new	1	1	0	1
2	new	1	1	0	1
3	accessible	1	0	0	1
4	accessible	1	0	0	1
..
648	accessible	1	1	0	1
649	accessible	1	1	0	1
650	accessible	0	0	0	1
651	accessible	1	1	0	1
652	accessible	0	0	0	1

[653 rows x 20 columns]

0.1.2 Part 2

Logistic regression model to the training set.

Predictors: recipient pronominality, intercept.

```
[ ]: X = train_set[["RecPro", "Dummy"]]
      y = train_set[["Response"]]

      model_1 = sm.GLM(y, X, family=sm.families.Binomial())
      model_1 = model_1.fit()
      print(model_1.summary())
```

```

                        Generalized Linear Model Regression Results
=====
Dep. Variable:                Response    No. Observations:                2610
Model:                        GLM        Df Residuals:                  2608
Model Family:                 Binomial   Df Model:                        1
Link Function:                 Logit     Scale:                          1.0000
Method:                       IRLS     Log-Likelihood:                   -1246.3
Date:                         Sat, 13 May 2023    Deviance:                         2492.7
Time:                         19:21:49    Pearson chi2:                     2.61e+03
No. Iterations:                5         Pseudo R-squ. (CS):                0.1668
Covariance Type:               nonrobust
=====

```

	coef	std err	z	P> z	[0.025	0.975]
RecPro	2.0750	0.102	20.406	0.000	1.876	2.274
Dummy	0.0142	0.064	0.223	0.824	-0.111	0.139

```
=====
```

```
[ ]: eval_model(model_1, test_set, predictors=["RecPro", "Dummy"])
```

model accuracy on the test set: 0.723

model log-likelihood on the test set: -436.818

First, we can see that the model's accuracy on the test set is greater than 0.5. Thus, this model has some degree of predictive power, since it's better than a random classifier.

Second, it appears that $\beta_{RecPro} = 2.075$. Hence we can conclude the following:

- Since $\beta_{RecPro} > 0$, the probability of the dative alternation being DO is higher when the recipient (in the sentence) is a pronoun ($x_{RecPro} = 1$).
- The odds-ratio is $e^{2.075}$ times larger, when the recipient (in the sentence) is a pronoun ($x_{RecPro} = 1$).

0.1.3 Part 3

Logistic regression model to the training set.

Predictors: recipient pronominality, theme pronominality, intercept.


```
[ ]: X = train_set[["RecPro", "ThemePro", "Dummy"]]
      y = train_set[["Response"]]

      model_2 = sm.GLM(y, X, family=sm.families.Binomial())
      model_2 = model_2.fit()
      print(model_2.summary())
```

```

                        Generalized Linear Model Regression Results
=====
Dep. Variable:          Response    No. Observations:          2610
Model:                  GLM        Df Residuals:              2607
Model Family:          Binomial    Df Model:                  2
Link Function:         Logit       Scale:                   1.0000
Method:                IRLS       Log-Likelihood:         -1046.7
Date:                  Sat, 13 May 2023    Deviance:              2093.4
Time:                  19:21:49    Pearson chi2:          2.58e+03
No. Iterations:        6          Pseudo R-squ. (CS):      0.2850
Covariance Type:       nonrobust
=====

```

	coef	std err	z	P> z	[0.025	0.975]
RecPro	2.8476	0.140	20.371	0.000	2.574	3.122
ThemePro	-3.0000	0.168	-17.893	0.000	-3.329	-2.671
Dummy	0.2297	0.067	3.421	0.001	0.098	0.361

```
=====
```

```
[ ]: eval_model(model_2, test_set, predictors=["RecPro", "ThemePro", "Dummy"])
```

model accuracy on the test set: 0.75
model log-likelihood on the test set: -379.969

First, we can see that both the model's accuracy and log-likelihood on the test set are greater than those of the previous one. Thus, we can conclude that this model has a greater predictive power (by these metrics).

Second, it appears that $\beta_{RecPro} = 2.848$. Hence we can conclude the following:

- Since $\beta_{RecPro} > 0$, the probability of the dative alternation being DO is higher when the recipient (in the sentence) is a pronoun
 $(x_{RecPro} = 1)$.
- The odds-ratio is $e^{2.848}$ times larger, when the recipient (in the sentence) is a pronoun
 $(x_{RecPro} = 1)$.

Third, it appears that $\beta_{ThemePro} = -3.0$. Hence we can conclude the following:

- Since $\beta_{ThemePro} < 0$, the probability of the dative alternation being DO is lower when the theme (in the sentence) is a pronoun
 $(x_{ThemePro} = 1)$.

- The odds-ratio is e^{-3} times smaller, when the recipient (in the sentence) is a pronoun ($x_{ThemePro} = 1$).

Moreover, we observe that $|\beta_{ThemePro}| > |\beta_{RecPro}|$ and $\beta_{ThemePro} < 0$. As a result, the probability of a sentence with both recipient and theme as pronouns being classified as DO, is lower than that of a sentence with none as pronouns (according to the model).

0.1.4 Part 4

In this part, we will test two more logostic regression models.

Part 4.1 Logistic regression model to the training set.

Predictors: recipient pronominality, theme pronominality, recipient length, theme length, intercept.

```
[ ]: X = train_set[["RecPro", "ThemePro", "LengthOfRecipient", "LengthOfTheme", "
    ↳"Dummy"]]
y = train_set[["Response"]]

model_3 = sm.GLM(y, X, family=sm.families.Binomial())
model_3 = model_3.fit()
print(model_3.summary())
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:                Response    No. Observations:                2610
Model:                        GLM         Df Residuals:                  2605
Model Family:                 Binomial    Df Model:                      4
Link Function:                 Logit       Scale:                        1.0000
Method:                       IRLS       Log-Likelihood:                 -896.22
Date:                         Sat, 13 May 2023    Deviance:                      1792.4
Time:                         19:21:49         Pearson chi2:                   3.45e+03
No. Iterations:                6             Pseudo R-squ. (CS):             0.3629
Covariance Type:               nonrobust
=====
```

```
=====
                                coef      std err          z      P>|z|      [0.025
-----
0.975]
-----
RecPro                2.4551        0.160     15.299      0.000        2.141
2.770
ThemePro             -2.7656        0.172    -16.103      0.000       -3.102
-2.429
LengthOfRecipient    -0.4065        0.042    -9.622      0.000       -0.489
-0.324
LengthOfTheme         0.2333        0.023     10.243      0.000         0.189
0.278
Dummy                 0.3342        0.152      2.193      0.028         0.036
```

0.633

=====

```
[ ]: eval_model(model_3, test_set, predictors=["RecPro", "ThemePro",  
      ↪ "LengthOfRecipient", "LengthOfTheme", "Dummy"])
```

model accuracy on the test set: 0.864

model log-likelihood on the test set: -315.059

Part 4.2 Logistic regression model to the training set.

Predictors: recipient pronominality, theme pronominality, log recipient length, log theme length, intercept.

```
[ ]: # we use log with base 2 according to the course's convention  
train_set["LogLengthOfRecipient"] = [math.log2(x) for x in  
      ↪ train_set["LengthOfRecipient"]]  
train_set["LogLengthOfTheme"] = [math.log2(x) for x in  
      ↪ train_set["LengthOfTheme"]]  
  
test_set["LogLengthOfRecipient"] = [math.log2(x) for x in  
      ↪ test_set["LengthOfRecipient"]]  
test_set["LogLengthOfTheme"] = [math.log2(x) for x in test_set["LengthOfTheme"]]  
  
X = train_set[["RecPro", "ThemePro", "LogLengthOfRecipient",  
      ↪ "LogLengthOfTheme", "Dummy"]]  
y = train_set[["Response"]]  
  
model_4 = sm.GLM(y, X, family=sm.families.Binomial())  
model_4 = model_4.fit()  
print(model_4.summary())
```

Generalized Linear Model Regression Results

```
=====
```

Dep. Variable:	Response	No. Observations:	2610
Model:	GLM	Df Residuals:	2605
Model Family:	Binomial	Df Model:	4
Link Function:	Logit	Scale:	1.0000
Method:	IRLS	Log-Likelihood:	-893.29
Date:	Sat, 13 May 2023	Deviance:	1786.6
Time:	19:21:49	Pearson chi2:	2.74e+03
No. Iterations:	6	Pseudo R-squ. (CS):	0.3643
Covariance Type:	nonrobust		

```
=====
```

```
=====
```

	coef	std err	z	P> z	[0.025
--	------	---------	---	------	--------

0.975]

RecPro	2.0772	0.180	11.571	0.000	1.725
2.429					
ThemePro	-2.5591	0.175	-14.616	0.000	-2.902
-2.216					
LogLengthOfRecipient	-0.9592	0.087	-11.072	0.000	-1.129
-0.789					
LogLengthOfTheme	0.6892	0.059	11.702	0.000	0.574
0.805					
Dummy	0.1922	0.159	1.211	0.226	-0.119
0.503					

```
[ ]: eval_model(model_4, test_set, predictors=["RecPro", "ThemePro",
↪ "LogLengthOfRecipient", "LogLengthOfTheme", "Dummy"])
```

model accuracy on the test set: 0.861
model log-likelihood on the test set: -312.804

TODO: interpretation goes here.

It appears that the first model has a higher accuracy on the test set than the second one (0.864 compared to 0.861 respectively). However, the first has a lower log-likelihood on the test set than the second one (-315.059 compared to -312.804 respectively).

Thus, according to the first metric (accuracy), the first model is better. On the other hand, according to the second metric (log-likelihood) the second model is the better one.

Therefore, we will choose a model according to the metric which represents our objective. However, notice that both metrics are very similar for the two models. Thus, both models will give us similar results.

0.1.5 Part 5

For the theoretical discussion of the regression coefficients, we will interpret the coefficients of the first model (non-log model).

1. It appears that $\beta_{RecPro} = 2.455$. Hence we can conclude the following:
 - Since $\beta_{RecPro} > 0$, the probability of the dative alternation being DO is higher when the recipient (in the sentence) is a pronoun ($x_{RecPro} = 1$).
 - The odds-ratio is $e^{2.455}$ times larger, when the recipient (in the sentence) is a pronoun ($x_{RecPro} = 1$).
2. It appears that $\beta_{ThemePro} = -2.766$. Hence we can conclude the following:
 - Since $\beta_{ThemePro} < 0$, the probability of the dative alternation being DO is lower when the theme (in the sentence) is a pronoun

$(x_{ThemePro} = 1)$.

- The odds-ratio is $e^{-2.766}$ times smaller, when the recipient (in the sentence) is a pronoun ($x_{ThemePro} = 1$).
3. It appears that $\beta_{RecLen} = -0.406$. Hence we can conclude the following:
- Since $\beta_{RecLen} < 0$, the probability of the dative alternation being DO decreases as the length of the recipient (in the sentence) increases.
 - The odds-ratio is $e^{-0.406}$ times smaller, when the length of the recipient (in the sentence) increases by one.
4. It appears that $\beta_{ThemeLen} = 0.334$. Hence we can conclude the following:
- Since $\beta_{ThemeLen} > 0$, the probability of the dative alternation being DO increases as the length of the theme (in the sentence) increases.
 - The odds-ratio is $e^{0.334}$ times larger, when the length of the theme (in the sentence) increases by one.

(*) As mentioned in the Pset PDF, we don't need to discuss the intercept (since its value depends on the numeric coding scheme used for the predictors).

Simplifying The Model Now, we will create a simplified model - which contains less explaining variables than the non-log full model.

In order to simplify the model and reduce the number of variables, we decided to combine the variables "LengthOfRecipient" and "LengthOfTheme" into a new variable "SqrtLengthSum", as follows:

$$SqrtLengthSum = (LengthOfRecipient + LengthOfTheme)^{\frac{1}{2}}$$

(Then, we removed the variables "LengthOfRecipient" and "LengthOfTheme").

Why we decided to create this variable:

Even though we lose information on the individual lengths of the recipient and the theme, the sum still holds some information regarding their lengths. In addition, we decided to take the square root of the sum, in order to scale it down.

```
[ ]: train_set["SqrtLengthSum"] = np.sqrt(train_set["LengthOfRecipient"] +  
    ↪ train_set["LengthOfTheme"])  
test_set["SqrtLengthSum"] = np.sqrt(test_set["LengthOfRecipient"] +  
    ↪ test_set["LengthOfTheme"])  
  
X = train_set[["RecPro", "ThemePro", "SqrtLengthSum", "Dummy"]]  
y = train_set[["Response"]]  
  
model_5 = sm.GLM(y, X, family=sm.families.Binomial())  
model_5 = model_5.fit()  
print(model_5.summary())
```

Generalized Linear Model Regression Results

=====

```

Dep. Variable:          Response    No. Observations:          2610
Model:                  GLM         Df Residuals:              2606
Model Family:          Binomial    Df Model:                  3
Link Function:         Logit       Scale:                    1.0000
Method:                IRLS        Log-Likelihood:           -1030.7
Date:                  Sat, 13 May 2023    Deviance:                 2061.5
Time:                  19:21:49    Pearson chi2:             2.58e+03
No. Iterations:        6          Pseudo R-squ. (CS):       0.2937
Covariance Type:      nonrobust

```

```

=====
=
              coef      std err          z      P>|z|      [0.025
0.975]
-----
-
RecPro          3.1146      0.150      20.814      0.000      2.821
3.408
ThemePro       -2.8693      0.170     -16.881      0.000     -3.202
-2.536
SqrtLengthSum   0.3912      0.072       5.465      0.000      0.251
0.532
Dummy          -0.8537      0.208      -4.112      0.000     -1.261
-0.447
=====
=

```

```
[ ]: eval_model(model_5, test_set, predictors=["RecPro", "ThemePro",
↪ "SqrtLengthSum", "Dummy"])
```

```

model accuracy on the test set: 0.795
model log-likelihood on the test set: -374.076

```

As we can see, the new model's accuracy is 0.795, compared to 0.864 by the full non-log model. Thus, the predictive accuracy sacrificed is $0.069 < 0.1$. Therefore, the accuracy lost is considered “not much” (according to the forums), as required.

Improving Our Solution We will now try another model in order to improve the accuracy of the previous suggested model.

Here, we decided to combine the variables “RecPro” and “ThemePro” into a new variable “SubPro”, as follows:

$$SubPro = RecPro - ThemePro$$

(Then, we removed the variables “RecPro” and “ThemePro”).

Why we decided to create this variable:

Notice that the values of the new variable ‘SubPro’ are: (-1, 0, 1).

Moreover, the following holds:

- $SubPro = 1$ IFF $RecPro = 1$ and $ThemePro = 0$.
- $SubPro = -1$ IFF $RecPro = 0$ and $ThemePro = 1$.
- $SubPro = 0$ IFF ($RecPro = ThemePro = 0$ or $RecPro = ThemePro = 1$).

We can see that information loss happens only when $SubPro = 0$ (two different cases are represented by a single value).

However, In the full non-log model, it appears that $\beta_{RecPro} = 2.455$ and $\beta_{ThemePro} = -2.766$.

Thus, we can see that $\beta_{RecPro} \approx -\beta_{ThemePro}$.

As a result, the probability of a sentence with both recipient and theme as pronouns being classified as DO, is approximately the same as the probability of a sentence with none as pronouns being classified as DO. In other words, the case where $RecPro = ThemePro = 0$ is roughly the same as the case where $RecPro = ThemePro = 1$ (according to the model).

Therefore, the new variable ‘SubPro’ successfully encodes both variables ‘RecPro’ and ‘ThemePro’ without losing a significant amount of information.

```
[ ]: train_set["SubPro"] = train_set["RecPro"] - train_set["ThemePro"]
test_set["SubPro"] = test_set["RecPro"] - test_set["ThemePro"]

X = train_set[["SubPro", "LengthOfRecipient", "LengthOfTheme", "Dummy"]]
y = train_set[["Response"]]

model_5_improved = sm.GLM(y, X, family=sm.families.Binomial())
model_5_improved = model_5_improved.fit()
print(model_5_improved.summary())
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:          Response    No. Observations:          2610
Model:                  GLM        Df Residuals:              2606
Model Family:           Binomial   Df Model:                  3
Link Function:           Logit     Scale:                   1.0000
Method:                 IRLS      Log-Likelihood:         -897.91
Date:                   Sat, 13 May 2023    Deviance:               1795.8
Time:                   19:21:49    Pearson chi2:           3.51e+03
No. Iterations:         6          Pseudo R-squ. (CS):      0.3620
Covariance Type:        nonrobust
=====
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
SubPro          2.5897      0.142     18.201      0.000      2.311
2.869
LengthOfRecipient -0.3800      0.039    -9.823      0.000     -0.456
-0.304
```

LengthOfTheme	0.2416	0.022	10.764	0.000	0.198
0.286					
Dummy	0.1593	0.117	1.356	0.175	-0.071
0.389					

```
=====
=====
```

```
[ ]: eval_model(model_5_improved, test_set, predictors=["SubPro",
↪ "LengthOfRecipient", "LengthOfTheme", "Dummy"])
```

model accuracy on the test set: 0.85

model log-likelihood on the test set: -314.611

As we can see, the new model's accuracy is 0.85, compared to 0.864 by the full non-log model.

Thus, the predictive accuracy sacrificed is 0.014 (roughly 5 times less than the previous suggested model).

Therefore, we succeeded in our task to improve our simplified model.